

City Prediction from Weather Data

July 9, 2019

1 Introduction

We have a weather dataset of 26 North American cities from roughly 1960 to 2015, downloaded from the Global Climatology Network. The dataset consists of the following monthly weather data for each city for each year: 1. average maximum temperature(0.1°C) 2. average minimum temperature(0.1°C) 3. average precipitation(0.1mm) 4. average snowfall(mm) 5. average snow depth(mm)

Using Python and relevant data science libraries, We plan to build and train an accurate machine learning algorithm to classify given unlabelled weather data into their corresponding North American cities. We will try a few different algorithms, validate them, and then choose the one with the best results.

2 Set-up

We need to import Python's data science libraries and read the data file.

```
In [8]: import numpy as np
import pandas as pd
```

```
data = pd.read_csv('monthly-data-labelled.csv')
#This is what the data looks like
data.head()
```

```
Out[8]:
```

	city	year	tmax-01	tmax-02	tmax-03	tmax-04	tmax-05	\
0	Anchorage	1960	-46.516129	-9.482759	-9.677419	52.400000	140.967742	
1	Anchorage	1961	-26.096774	-44.571429	-35.064516	58.200000	140.193548	
2	Anchorage	1962	-59.225806	-31.750000	-18.903226	69.366667	111.419355	
3	Anchorage	1963	-39.290323	-11.357143	-1.451613	41.700000	134.258065	
4	Anchorage	1964	-59.129032	-24.655172	-35.096774	45.866667	99.903226	

	tmax-06	tmax-07	tmax-08	...	snwd-03	snwd-04	snwd-05	\
0	173.166667	180.225806	168.064516	...	290.903226	44.066667	0.0	
1	169.633333	178.645161	161.806452	...	113.096774	8.433333	0.0	
2	159.633333	187.451613	176.483871	...	128.645161	5.866667	0.0	
3	146.200000	185.612903	182.129032	...	60.645161	78.766667	0.0	
4	173.566667	182.516129	163.483871	...	114.793103	53.266667	0.0	

	snwd-06	snwd-07	snwd-08	snwd-09	snwd-10	snwd-11	snwd-12
0	0.0	0.0	0.0	0.0	0.000000	29.433333	77.612903
1	0.0	0.0	0.0	0.0	45.032258	98.366667	147.258065
2	0.0	0.0	0.0	0.0	0.000000	10.000000	46.548387
3	0.0	0.0	0.0	0.0	8.161290	34.466667	18.032258
4	0.0	0.0	0.0	0.0	15.516129	148.133333	345.870968

[5 rows x 62 columns]

From this dataframe, we need a two-dimensional array consisting of weather data values without the city label, and another array consisting of the corresponding city labels. Machine learning models are trained using this format of dataset.

```
In [2]: X = data.drop('city', 1).values #remove 'city' column
        y = data['city'].values
```

3 Training and Validating

We will try three types of machine learning models: Naive Bayesian, k-Nearest Neighbours, and Support Vector Machine. We will create pipelines to scale each feature's values to lie between 0 and 1 before computing the classifications.

For each model, we will divide the data into training and validation sets. This is how we will find out the accuracy of the trained model - by testing the model on the validation set, which will have never been shown to the model during training.

```
In [3]: from sklearn.model_selection import train_test_split
```

3.1 Naive Bayes (NB)

```
In [4]: from sklearn.pipeline import make_pipeline
        from sklearn.preprocessing import MinMaxScaler
        from sklearn.naive_bayes import GaussianNB

        n = 30 #number of iterations
        total = 0
        for i in range(n):
            X_train, X_valid, y_train, y_valid = train_test_split(X, y)
            bayes_model = make_pipeline(
                MinMaxScaler(),
                GaussianNB()
            )
            bayes_model.fit(X_train, y_train)
            total += bayes_model.score(X_valid, y_valid)

        print("bayes model score: " + str(total/n))
```

bayes model score: 0.6514942528735633

The bayes model has a score of roughly 65% accuracy, which is terrible. Why is this the case? NB algorithm assumes input variables to be independent, and also assumes input data to be distributed normally. None of these assumptions would hold true for weather data. For example, temperature is highly correlated with snowfall and snow depth. Moreover, monthly temperatures throughout the decades have no reason to be normally distributed. This explains the model's low score.

3.2 k-Nearest Neighbours (kNN)

```
In [5]: from sklearn.neighbors import KNeighborsClassifier
```

```
n_nbs = 5 #number of neighbours to consider
n = 30 #number of iterations
total = 0
for i in range(n):
    X_train, X_valid, y_train, y_valid = train_test_split(X, y)
    knn_model = make_pipeline(
        MinMaxScaler(),
        KNeighborsClassifier(n_neighbors=n_nbs)
    )
    knn_model.fit(X_train, y_train)
    total += knn_model.score(X_valid, y_valid)

print("knn model score: " + str(total/n))
```

```
knn model score: 0.6647126436781612
```

The result is around 65%, which is as bad as the NB approach. One guess as to why this algorithm is not effective is because we have mapped the values of all the features to the same scale. Since kNN measures the Euclidian distance between data points in the feature space to determine the class of the prediction input, scaling as mentioned results in all the features contributing approximately the same weight to the distance calculation. We have no knowledge of which features most significantly affect classification, and it could be the case that some features should be more heavily weighed than others.

3.3 Support Vector Machine (SVM)

```
In [6]: from sklearn.svm import SVC
```

```
n = 30 #number of iterations
total = 0
for i in range(n):
    X_train, X_valid, y_train, y_valid = train_test_split(X, y)
    svc_model = make_pipeline(
        MinMaxScaler(),
        SVC(kernel='rbf', C=5, gamma='scale', decision_function_shape='ovr')
    )
```

```
svc_model.fit(X_train, y_train)
total += svc_model.score(X_valid, y_valid)

print("svc model score: " + str(svc_model.score(X_valid, y_valid)))

svc model score: 0.8137931034482758
```

Our SVM model's accuracy is around 80% - much better than the other two. SVM draws optimal hyperplanes between classes, which is more flexible and accurate than the methods of the other two machine learning algorithms.

4 Conclusion

We trained three machine learning models (Naive Bayes, k-Nearest Neighbours, and Support Vector Machine) using a relatively small decades-long weather dataset with 60 features (5 weather measurements * 12 months). The models were trained to predict cities from weather data, then were evaluated for their accuracy. The SVM model had the top accuracy of around 80%, while the other two models had an accuracy of around 66%.

Our SVM model is far from being practically useful. The model can be improved by adding more rows; cities like Denver has fewer data points than the rest of the cities. Perhaps Denver's (and other cities') missing weather information can be researched and added to the table. Another way to improve the model is to add more relevant features to our data. Adding wind velocity and atmospheric pressure, for example, will significantly improve all our models' accuracy.

In []: