

City Prediction from Weather Data

July 7, 2019

1 Introduction

We have a weather dataset of 26 North American cities from roughly 1960 to 2015, downloaded from the Global Climatology Network (<https://www.ncdc.noaa.gov/data-access/land-based-station-data/land-based-datasets/global-historical-climatology-network-ghcn>). The dataset consists of the following monthly weather data for each city: 1. average maximum temperature(0.1°C) 2. average minimum temperature(0.1°C) 3. average precipitation(0.1mm) 4. average snow-fall(mm) 5. average snow depth(mm)

Using Python and relevant data science libraries, We plan to build and train an accurate machine learning algorithm to classify given unlabelled weather data into their corresponding North American cities. We will try a few different algorithms, validate them, and then choose the one with the best results.

2 Set-up

We need to import Python's data science libraries and read the data file.

```
In [1]: import numpy as np
import pandas as pd
```

```
data = pd.read_csv('monthly-data-labelled.csv')
#This is what the data looks like
data
```

```
Out[1]:
```

	city	year	tmax-01	tmax-02	tmax-03	tmax-04	\
0	Anchorage	1960	-46.516129	-9.482759	-9.677419	52.400000	
1	Anchorage	1961	-26.096774	-44.571429	-35.064516	58.200000	
2	Anchorage	1962	-59.225806	-31.750000	-18.903226	69.366667	
3	Anchorage	1963	-39.290323	-11.357143	-1.451613	41.700000	
4	Anchorage	1964	-59.129032	-24.655172	-35.096774	45.866667	
5	Anchorage	1965	-81.193548	-73.607143	53.387097	84.266667	
6	Anchorage	1966	-82.387097	-47.285714	-36.419355	57.266667	
7	Anchorage	1967	-93.741935	-50.821429	1.290323	55.266667	
8	Anchorage	1968	-75.580645	-20.310345	17.000000	49.066667	
9	Anchorage	1969	-112.322581	-35.285714	15.677419	79.166667	
10	Anchorage	1970	-87.967742	15.107143	50.354839	59.233333	

11	Anchorage	1971	-121.838710	-27.750000	-47.225806	41.366667
12	Anchorage	1972	-104.677419	-55.517241	-47.161290	10.400000
13	Anchorage	1973	-119.419355	-58.571429	-6.354839	51.766667
14	Anchorage	1974	-99.838710	-65.142857	1.322581	74.600000
15	Anchorage	1975	-69.548387	-62.642857	-9.483871	35.566667
16	Anchorage	1976	-54.193548	-62.241379	-15.032258	51.033333
17	Anchorage	1977	24.354839	35.750000	-2.129032	53.933333
18	Anchorage	1978	-28.419355	5.214286	24.580645	83.700000
19	Anchorage	1979	-19.967742	-64.035714	30.322581	71.966667
20	Anchorage	1980	-61.677419	9.586207	11.290323	86.900000
21	Anchorage	1981	28.161290	-10.357143	51.935484	72.600000
22	Anchorage	1982	-100.161290	-54.928571	4.322581	47.200000
23	Anchorage	1983	-63.516129	-23.500000	26.258065	66.000000
24	Anchorage	1984	-38.354839	-34.482759	65.322581	78.466667
25	Anchorage	1985	22.645161	-62.928571	12.548387	23.666667
26	Anchorage	1986	-7.161290	-18.785714	2.548387	36.500000
27	Anchorage	1987	-17.548387	0.321429	16.290323	76.866667
28	Anchorage	1988	-44.516129	-13.965517	27.258065	65.200000
29	Anchorage	1989	-116.322581	-37.928571	3.064516	83.100000
...
1129	Winnipeg	1973	-89.000000	-84.142857	46.741935	91.400000
1130	Winnipeg	1974	-170.741935	-111.464286	-50.483871	69.566667
1131	Winnipeg	1975	-108.548387	-108.321429	-56.225806	68.133333
1132	Winnipeg	1976	-121.387097	-57.965517	-39.096774	127.666667
1133	Winnipeg	1977	-160.580645	-65.785714	22.741935	160.133333
1134	Winnipeg	1978	-168.354839	-122.642857	-27.354839	87.133333
1135	Winnipeg	1979	-178.451613	-167.535714	-50.870968	38.800000
1136	Winnipeg	1980	-139.096774	-111.551724	-40.258065	161.000000
1137	Winnipeg	1981	-84.322581	-54.857143	44.741935	118.000000
1138	Winnipeg	1982	-198.806452	-95.321429	-24.645161	91.100000
1139	Winnipeg	1983	-77.935484	-60.750000	-12.322581	90.933333
1140	Winnipeg	1984	-108.483871	-24.586207	-18.709677	141.933333
1141	Winnipeg	1985	-126.032258	-114.321429	18.354839	138.633333
1142	Winnipeg	1986	-80.000000	-103.071429	8.451613	101.266667
1143	Winnipeg	1987	-70.290323	-25.571429	-1.806452	161.433333
1144	Winnipeg	1988	-131.096774	-98.931034	5.612903	135.433333
1145	Winnipeg	1989	-100.451613	-131.964286	-55.129032	82.233333
1146	Winnipeg	1990	-65.741935	-86.285714	3.645161	90.000000
1147	Winnipeg	1991	-131.354839	-54.035714	-5.225806	134.000000
1148	Winnipeg	1992	-67.967742	-53.655172	-1.096774	67.666667
1149	Winnipeg	1993	-95.526316	-94.238095	-19.807692	106.040000
1150	Winnipeg	1994	-186.741935	-130.892857	28.870968	105.800000
1151	Winnipeg	1995	-110.258065	-101.000000	-10.677419	65.500000
1152	Winnipeg	1996	-174.612903	-100.068966	-73.354839	30.833333
1153	Winnipeg	1997	-153.290323	-76.785714	-51.032258	41.766667
1154	Winnipeg	1998	-108.709677	-15.892857	-12.548387	155.833333
1155	Winnipeg	1999	-131.838710	-34.892857	27.193548	132.366667
1156	Winnipeg	2000	-114.774194	-36.344828	60.161290	118.300000

1157	Winnipeg	2001	-72.387097	-131.607143	-10.548387	109.600000
1158	Winnipeg	2002	-97.741935	-41.714286	-60.806452	81.833333

	tmax-05	tmax-06	tmax-07	tmax-08	...	snwd-03 \
0	140.967742	173.166667	180.225806	168.064516	...	290.903226
1	140.193548	169.633333	178.645161	161.806452	...	113.096774
2	111.419355	159.633333	187.451613	176.483871	...	128.645161
3	134.258065	146.200000	185.612903	182.129032	...	60.645161
4	99.903226	173.566667	182.516129	163.483871	...	114.793103
5	121.322581	150.200000	188.161290	176.870968	...	227.741935
6	105.612903	170.966667	183.612903	153.483871	...	476.967742
7	136.774194	166.766667	190.580645	176.483871	...	285.935484
8	132.935484	170.233333	189.967742	187.580645	...	40.806452
9	128.322581	184.066667	185.064516	167.548387	...	253.322581
10	134.967742	159.266667	170.483871	157.354839	...	27.806452
11	85.258065	145.933333	165.516129	161.774194	...	241.838710
12	95.838710	145.666667	196.387097	172.774194	...	239.225806
13	101.129032	146.566667	181.000000	153.322581	...	352.322581
14	144.064516	175.566667	181.096774	180.709677	...	243.290323
15	120.870968	155.466667	183.419355	170.129032	...	386.741935
16	110.161290	164.966667	189.387097	170.935484	...	174.483871
17	124.032258	187.200000	213.354839	195.000000	...	49.870968
18	139.225806	160.733333	185.064516	203.225806	...	323.677419
19	145.193548	177.600000	192.064516	185.645161	...	379.354839
20	113.000000	152.233333	174.612903	161.548387	...	24.483871
21	152.870968	162.533333	169.903226	157.935484	...	35.290323
22	114.645161	155.200000	169.935484	168.225806	...	46.580645
23	133.096774	179.700000	185.322581	171.419355	...	33.483871
24	146.258065	190.733333	192.483871	175.258065	...	274.677419
25	116.193548	149.733333	182.064516	165.612903	...	261.290323
26	126.516129	170.200000	184.806452	155.612903	...	88.451613
27	123.322581	144.366667	172.516129	184.516129	...	99.903226
28	131.451613	166.133333	183.387097	168.483871	...	336.903226
29	121.483871	174.300000	193.870968	183.451613	...	132.032258
...
1129	183.290323	219.400000	244.129032	266.709677	...	27.096774
1130	135.870968	246.200000	289.483871	226.516129	...	271.935484
1131	190.064516	229.566667	284.387097	222.580645	...	234.193548
1132	201.935484	241.700000	266.806452	263.258065	...	269.032258
1133	244.032258	220.333333	256.548387	207.903226	...	106.129032
1134	213.580645	227.366667	248.064516	241.548387	...	133.870968
1135	126.387097	227.066667	274.548387	228.709677	...	317.419355
1136	238.032258	240.466667	273.387097	225.419355	...	201.034483
1137	187.645161	223.000000	268.322581	265.903226	...	1.290323
1138	196.129032	205.900000	259.548387	230.387097	...	16.451613
1139	159.161290	233.433333	280.580645	296.806452	...	34.193548
1140	175.225806	223.833333	258.935484	282.387097	...	20.645161
1141	206.354839	198.600000	257.064516	214.870968	...	101.290323

1142	206.032258	231.433333	249.000000	244.064516	...	87.741935
1143	225.709677	261.500000	257.000000	235.258065	...	199.677419
1144	224.096774	294.400000	279.741935	270.419355	...	18.064516
1145	214.161290	228.833333	283.225806	260.774194	...	160.645161
1146	180.161290	238.466667	254.935484	274.806452	...	31.612903
1147	202.451613	242.433333	250.967742	279.709677	...	85.483871
1148	199.387097	207.066667	211.225806	221.967742	...	39.677419
1149	182.227273	208.600000	222.612903	229.777778	...	136.538462
1150	197.612903	237.700000	239.806452	232.419355	...	14.193548
1151	178.516129	270.533333	259.193548	268.193548	...	96.451613
1152	161.193548	251.366667	251.774194	261.193548	...	431.200000
1153	154.000000	260.400000	251.516129	245.612903	...	448.387097
1154	198.677419	205.000000	259.774194	279.806452	...	60.645161
1155	177.870968	222.233333	254.741935	248.096774	...	40.645161
1156	186.903226	200.400000	257.451613	256.451613	...	1.290323
1157	189.612903	223.266667	263.516129	259.580645	...	157.096774
1158	155.645161	243.166667	274.741935	246.161290	...	104.838710

	snwd-04	snwd-05	snwd-06	snwd-07	snwd-08	snwd-09	snwd-10 \
0	44.066667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1	8.433333	0.000000	0.0	0.0	0.0	0.000000	45.032258
2	5.866667	0.000000	0.0	0.0	0.0	0.000000	0.000000
3	78.766667	0.000000	0.0	0.0	0.0	0.000000	8.161290
4	53.266667	0.000000	0.0	0.0	0.0	0.000000	15.516129
5	1.666667	0.000000	0.0	0.0	0.0	0.833333	53.161290
6	102.400000	0.000000	0.0	0.0	0.0	0.000000	0.806452
7	35.533333	0.000000	0.0	0.0	0.0	0.000000	13.161290
8	20.233333	0.000000	0.0	0.0	0.0	0.000000	63.096774
9	48.233333	0.000000	0.0	0.0	0.0	0.000000	0.000000
10	10.966667	0.000000	0.0	0.0	0.0	0.000000	8.096774
11	105.766667	0.000000	0.0	0.0	0.0	0.000000	36.032258
12	292.066667	7.322581	0.0	0.0	0.0	0.000000	3.258065
13	63.500000	0.000000	0.0	0.0	0.0	0.000000	8.129032
14	0.833333	0.000000	0.0	0.0	0.0	0.000000	5.677419
15	260.733333	2.451613	0.0	0.0	0.0	0.000000	0.000000
16	88.066667	0.000000	0.0	0.0	0.0	0.000000	42.580645
17	35.533333	0.000000	0.0	0.0	0.0	0.000000	30.290323
18	47.433333	0.000000	0.0	0.0	0.0	0.000000	1.645161
19	89.733333	0.000000	0.0	0.0	0.0	0.000000	1.612903
20	0.000000	0.000000	0.0	0.0	0.0	0.000000	28.709677
21	1.666667	0.000000	0.0	0.0	0.0	0.000000	2.419355
22	2.500000	0.000000	0.0	0.0	0.0	0.000000	79.419355
23	5.900000	0.000000	0.0	0.0	0.0	0.000000	36.870968
24	14.300000	0.000000	0.0	0.0	0.0	0.000000	4.096774
25	213.400000	0.000000	0.0	0.0	0.0	0.000000	0.000000
26	106.600000	0.000000	0.0	0.0	0.0	0.000000	0.000000
27	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
28	134.666667	0.000000	0.0	0.0	0.0	0.000000	10.612903

29	5.100000	0.000000	0.0	0.0	0.0	0.000000	60.645161
...
1129	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1130	130.000000	0.000000	0.0	0.0	0.0	0.000000	0.967742
1131	51.666667	0.967742	0.0	0.0	0.0	0.000000	0.967742
1132	4.333333	0.000000	0.0	0.0	0.0	0.000000	0.000000
1133	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1134	1.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1135	139.333333	1.290323	0.0	0.0	0.0	0.000000	0.000000
1136	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.967742
1137	0.333333	0.000000	0.0	0.0	0.0	0.000000	1.935484
1138	2.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1139	0.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1140	1.666667	0.000000	0.0	0.0	0.0	0.666667	3.870968
1141	0.333333	0.000000	0.0	0.0	0.0	0.000000	7.419355
1142	1.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1143	1.000000	0.000000	0.0	0.0	0.0	0.000000	2.903226
1144	0.000000	0.000000	0.0	0.0	0.0	0.000000	8.965517
1145	5.333333	0.000000	0.0	0.0	0.0	0.000000	0.000000
1146	1.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1147	5.000000	3.548387	0.0	0.0	0.0	0.000000	4.193548
1148	3.333333	0.000000	0.0	0.0	0.0	0.000000	0.000000
1149	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1150	3.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1151	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1152	208.666667	0.000000	0.0	0.0	0.0	0.000000	0.645161
1153	263.076923	0.322581	0.0	0.0	0.0	0.000000	0.000000
1154	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1155	22.666667	0.000000	0.0	0.0	0.0	0.000000	0.000000
1156	0.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1157	6.000000	0.000000	0.0	0.0	0.0	0.000000	0.000000
1158	0.333333	2.580645	0.0	0.0	0.0	0.000000	1.290323

	snwd-11	snwd-12
0	29.433333	77.612903
1	98.366667	147.258065
2	10.000000	46.548387
3	34.466667	18.032258
4	148.133333	345.870968
5	110.000000	458.774194
6	96.533333	287.580645
7	12.633333	234.354839
8	228.666667	364.774194
9	17.500000	34.322581
10	32.133333	111.580645
11	152.266667	240.129032
12	60.733333	242.483871
13	61.600000	137.612903

14	54.133333	241.677419
15	20.966667	85.193548
16	54.133333	176.096774
17	154.733333	176.967742
18	39.766667	360.354839
19	35.533333	95.774194
20	15.100000	26.677419
21	103.233333	172.806452
22	149.266667	210.516129
23	144.666667	208.225806
24	5.833333	88.612903
25	1.666667	1.612903
26	2.533333	42.548387
27	143.133333	490.806452
28	77.666667	182.677419
29	98.100000	153.193548
...
1129	94.666667	207.419355
1130	7.000000	44.193548
1131	7.000000	120.645161
1132	0.000000	74.516129
1133	49.333333	151.612903
1134	97.333333	232.258065
1135	27.000000	74.193548
1136	22.666667	24.193548
1137	1.666667	15.806452
1138	6.000000	22.258065
1139	22.333333	48.387097
1140	7.666667	79.032258
1141	58.000000	126.129032
1142	177.000000	131.612903
1143	0.000000	24.193548
1144	46.666667	85.161290
1145	15.333333	52.903226
1146	9.333333	81.290323
1147	49.333333	159.032258
1148	143.000000	253.333333
1149	16.000000	70.000000
1150	26.000000	75.517241
1151	120.333333	27.096774
1152	149.666667	409.354839
1153	26.666667	10.322581
1154	8.333333	66.774194
1155	4.333333	33.548387
1156	37.000000	208.064516
1157	0.666667	123.870968
1158	9.000000	45.483871

```
[1159 rows x 62 columns]
```

From this dataframe, we need a two-dimensional array consisting of weather data values without the city label, and another array consisting of the corresponding city labels. Machine learning models are trained using this format of dataset.

```
In [2]: X = data.drop('city', 1).values #remove 'city' column
        y = data['city'].values
```

3 Training and Validating

We will try three types of machine learning models: Naive Bayesian, k-Nearest Neighbours, and Support Vector Machine. We will create pipelines to scale each feature's values to lie between 0 and 1 before computing the classifications.

For each model, we will divide the data into training and validation sets. This is how we will find out the accuracy of the trained model - by testing the model on the validation set, which will have never been shown to the model during training.

```
In [3]: from sklearn.model_selection import train_test_split
```

3.1 Naive Bayesian (NB)

```
In [4]: from sklearn.pipeline import make_pipeline
        from sklearn.preprocessing import MinMaxScaler
        from sklearn.naive_bayes import GaussianNB

        n = 30 #number of iterations
        total = 0
        for i in range(n):
            X_train, X_valid, y_train, y_valid = train_test_split(X, y)
            bayes_model = make_pipeline(
                MinMaxScaler(),
                GaussianNB()
            )
            bayes_model.fit(X_train, y_train)
            total += bayes_model.score(X_valid, y_valid)

        print("bayes model score: " + str(total/n))
```

```
bayes model score: 0.6533333333333334
```

The bayes model has a score of roughly 65% accuracy, which is terrible. Why is this the case? NB algorithm assumes input variables to be independent, and also assumes input data to be distributed normally. None of these assumptions would hold true for weather data. For example, temperature is highly correlated with snowfall and snow depth. Moreover, monthly temperatures throughout the decades have no reason to be normally distributed. This explains the model's low score.

3.2 k-Nearest Neighbours (kNN)

```
In [5]: from sklearn.neighbors import KNeighborsClassifier

n_nbs = 5 #number of neighbours to consider
n = 30 #number of iterations
total = 0
for i in range(n):
    X_train, X_valid, y_train, y_valid = train_test_split(X, y)
    knn_model = make_pipeline(
        MinMaxScaler(),
        KNeighborsClassifier(n_neighbors=n_nbs)
    )
    knn_model.fit(X_train, y_train)
    total += knn_model.score(X_valid, y_valid)

print("knn model score: " + str(total/n))
```

knn model score: 0.6677011494252874

The result is as bad as the NB approach. One guess as to why this algorithm is not effective is because we have mapped the values of all the features to the same scale. Since kNN measures the Euclidian distance between data points in the feature space to determine the class of the prediction input, scaling as mentioned results in all the features contributing approximately the same weight to the distance calculation. We have no knowledge of which features most significantly affect classification, and it could be the case that some features should be more heavily weighed than others.

3.3 Support Vector Machine (SVM)

```
In [6]: from sklearn.svm import SVC

n = 30 #number of iterations
total = 0
for i in range(n):
    X_train, X_valid, y_train, y_valid = train_test_split(X, y)
    svc_model = make_pipeline(
        MinMaxScaler(),
        SVC(kernel='rbf', C=5, gamma='scale', decision_function_shape='ovr')
    )
    svc_model.fit(X_train, y_train)
    total += svc_model.score(X_valid, y_valid)

print("svc model score: " + str(svc_model.score(X_valid, y_valid)))
```

svc model score: 0.8482758620689655

Our SVM model's accuracy is much better than the other two. SVM is generally more robust to outliers/errors, since it considers all the outliers at the boundary between a pair of classes simultaneously. When more outliers are taken into account simultaneously, they tend to cancel out, which results in a good approximation of the boundary.

4 Conclusion

We trained three machine learning models (Naive Bayes, k-Nearest Neighbours, and Support Vector Machine) using a relatively small decades-long weather dataset with 60 features (5 weather measurements * 12 months). The models were trained to predict cities from weather data, then were evaluated for their accuracy. The SVM model had the top accuracy of around 80%, while the other two models had an accuracy of around 66%.

Our SVM model is far from being practically useful. The model can be improved by adding more rows; cities like Denver has fewer data points than the rest of the cities. Perhaps Denver's (and other cities') missing weather information can be researched and added to the table. Another way to improve the model is to add more relevant features to our data. Adding wind velocity and atmospheric pressure, for example, will significantly improve all our models' accuracy.

In []: